

Lecture 20

Memory Hierarchy

School of Computer Science and Engineering
Soongsil University

5. Large and Fast: Exploiting Memory Hierarchy

5.1 Introduction

5.2 Memory Technologies

5.3 The Basics of Caches

5.4 Measuring and Improving Cache Performance

5.5 Dependable Memory Hierarchy

5.6 Virtual Machines

5.7 Virtual Memory

5.8 A Common Framework for Memory Hierarchies

5.9 Using a Finite-State Machine to Control a Simple Cache

5.10 Parallelism and Memory Hierarchies: Cache Coherence


5.11 Parallelism and Memory Hierarchy: RAID

5.12 Advanced Material: Implementing Cache Controllers

5.13 Real Stuff: The ARM Cortex-A8 and Intel Core i7 Memory Hierarchies

5.1 Introduction


■ Processor-Memory Performance Gap



Apple II (1977): CPU...1000ns, Memory ... 400ns

The “Memory Wall”

- Processor vs. DRAM speed disparity continues to grow
- Good memory hierarchy (cache) design is increasingly important to overall performance



Memory Hierarchy

- **Engineering problem**

- ❖ Maximum performance with minimum cost

- **Fact**

- ❖ Fast memory → expensive → small
- ❖ Slow memory → inexpensive → large

- **Goal**

- ❖ To create the illusion of a large memory that we can access as fast as a very small memory

- **Multiple levels of memory**

- ❖ Faster memory: Close to the processor
- ❖ Slower and less expensive memory: Below that

- **Principle of locality**

- ❖ Programs access a small proportion of their address space at any time

Principle of Locality


1. Temporal locality

- ❖ Locality in time
 - ❖ Items accessed recently are likely to be accessed again soon
- [ex] loops

2. Spatial locality

- ❖ Locality in space
- ❖ Items near those accessed recently are likely to be accessed soon

[ex] sequential execution
array accesses



Basic Structure of a Memory Hierarchy

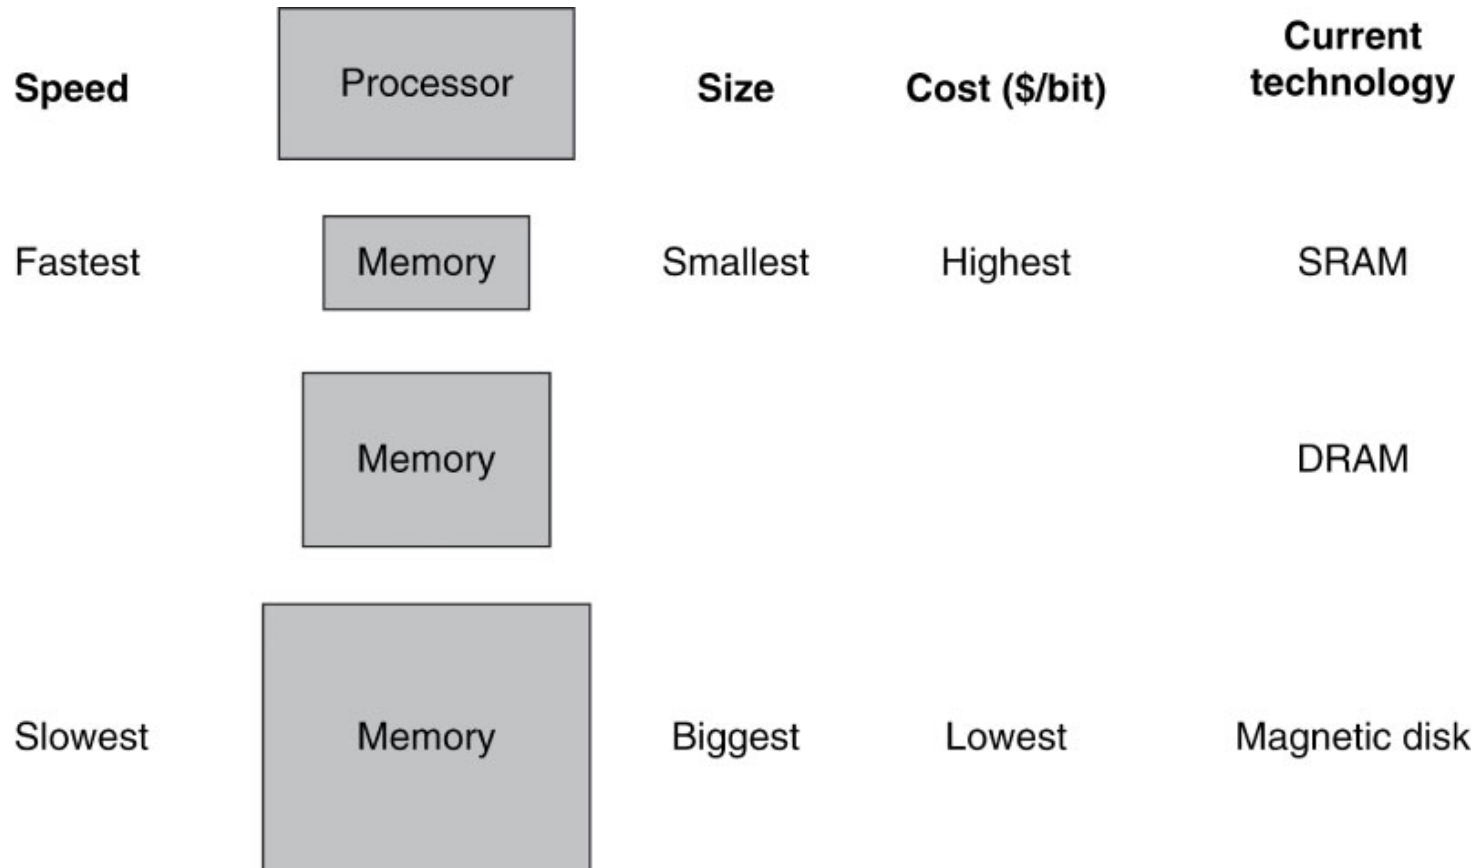
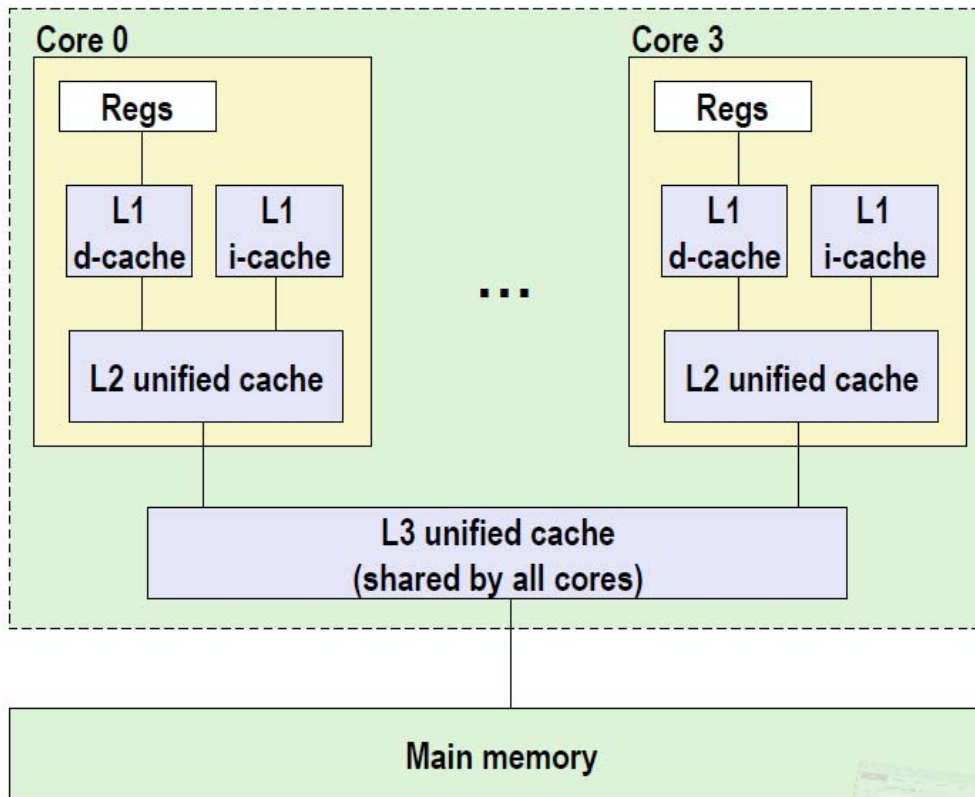


Figure 5.1

Memory Hierarchy in a Multicore Chip

Intel Core i7 Cache Hierarchy

Processor package



L1 i-cache and d-cache:

32 KB, 8-way,
Access: 4 cycles

L2 unified cache:

256 KB, 8-way,
Access: 11 cycles

L3 unified cache:

8 MB, 16-way,
Access: 30-40 cycles

Block size: 64 bytes for
all caches.

ARM Cortex-A8 and Intel Core i7 920

Characteristic	ARM Cortex-A8	Intel Core i7 920
L1 cache organization	Split instruction and data caches	Split instruction and data caches
L1 cache size	32 KiB each for instructions/data	32 KiB each for instructions/data per core
L1 cache associativity	4-way (I), 4-way (D) set associative	4-way (I), 8-way (D) set associative
L1 replacement	Random	Approximated LRU
L1 block size	64 bytes	64 bytes
L1 write policy	Write-back, Write-allocate(?)	Write-back, No-write-allocate
L1 hit time (load-use)	1 clock cycle	4 clock cycles, pipelined
L2 cache organization	Unified (instruction and data)	Unified (instruction and data) per core
L2 cache size	128 KiB to 1 MiB	256 KiB (0.25 MiB)
L2 cache associativity	8-way set associative	8-way set associative
L2 replacement	Random(?)	Approximated LRU
L2 block size	64 bytes	64 bytes
L2 write policy	Write-back, Write-allocate (?)	Write-back, Write-allocate
L2 hit time	11 clock cycles	10 clock cycles
L3 cache organization	-	Unified (instruction and data)
L3 cache size	-	8 MiB, shared
L3 cache associativity	-	16-way set associative
L3 replacement	-	Approximated LRU
L3 block size	-	64 bytes
L3 write policy	-	Write-back, Write-allocate
L3 hit time	-	35 clock cycles

Figure 5.44

Hit and Miss

- **Block (= line)**

- ❖ Minimum unit of information transfer between the two levels

- **Hit**

- ❖ When the data requested by the processor is in the upper level

- **Miss**

- ❖ When the data is not found in the upper level
- ❖ Then access the lower level

- **Hit rate (=hit ratio)**

- ❖ The fraction of memory accesses found in the upper level

- **Miss rate**

- ❖ $1 - \text{hit rate}$

Performance of the Memory Hierarchy

- **Hit time**

- ❖ Time to access the upper level of the memory hierarchy
- ❖ Time to access the block + Time to determine hit/miss

- **Miss penalty**

- ❖ Time to access the block in the lower level
 - + Time to transmit that block to the level that experienced the miss
 - + Time to insert the block in that level
 - + Time to pass the block to the requestor

- **Average memory access time (AMAT)**

- ❖ $\text{hit time} + \text{miss rate} \times \text{miss penalty}$

The BIG Picture

- **Temporal locality and memory hierarchy**
 - ❖ Keeping more recently accessed data items closer to the processor
- **Spatial locality and memory hierarchy**
 - ❖ Moving blocks consisting of multiple contiguous words to upper level
- **Memory hierarchy with high hit rate**
 - ❖ Access time: close to that of the highest level
 - ❖ Size: equal to that of the lowest level
- **Multi-level inclusion property**
 - ❖ $\text{Level}(i) \subset \text{Level}(i+1)$

5.2 Memory Technologies

Memory technology	Typical access time	\$ per GiB in 2012
SRAM semiconductor memory	0.5–2.5 ns	\$500–\$1000
DRAM semiconductor memory	50–70 ns	\$10–\$20
Flash semiconductor memory	5,000–50,000 ns	\$0.75–\$1.00
Magnetic disk	5,000,000–20,000,000 ns	\$0.05–\$0.10

- **SRAM (Static random access memory)**

- ❖ Fixed access time to any datum, though the read and write access times may differ
- ❖ No need to refresh, and so the access time is very close to the cycle time
- ❖ Six to eight transistors per bit

- **DRAM (Dynamic random access memory)**

- ❖ Value kept in a cell is stored as a charge in a capacitor ... 1 transistor/bit
- ❖ Dynamic: needs to be “refreshed” regularly (~ every 8 ms)
 - ◆ Consumes 1% to 2% of the active cycles of the DRAM

5.3 The Basics of Caches

- **Definition of Cache**

- ❖ The level of memory hierarchy between CPU and main memory
- ❖ Any storage managed to take advantage of locality of access

- **The first paper ... Wilkes[1965]**

- ❖ "Slave memories and dynamic storage allocation"

- **The first implementation**

- ❖ At the University of Cambridge by Scarrott

- **The first commercial machine with a cache**

- ❖ IBM 360/85, late 1960s

- **The first usage of the term "cache"**

- ❖ Conti, Gibson and Pitkowsky
- ❖ a paper in IBM Systems Journal in 1968

Simple Cache Scenario




Figure 5.7

■ 2 questions

- ❖ Q1: How do we know if a data item is in the cache?
- ❖ Q2: If it is, how do we find it?